

```
In [1]: # Imports for this Notebook
import pandas as pd
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn import metrics
from keras.models import Sequential
from keras.layers.core import Dense, Activation
from keras.callbacks import EarlyStopping
```

Using TensorFlow backend.

```
In [2]: from sklearn import preprocessing
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import shutil
import os
import keras
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: # Encode text values to dummy variables(i.e. [1,0,0],[0,1,0],[0,0,1] for r
def encode_text_dummy(df, name):
    dummies = pd.get_dummies(df[name])
    for x in dummies.columns:
        dummy_name = "{}-{}".format(name, x)
        df[dummy_name] = dummies[x]
    df.drop(name, axis=1, inplace=True)

# Encode text values to a single dummy variable. The new columns (which d
# at every location where the original column (name) matches each of the t
# each target value.
def encode_text_single_dummy(df, name, target_values):
    for tv in target_values:
        l = list(df[name].astype(str))
        l = [1 if str(x) == str(tv) else 0 for x in l]
        name2 = "{}-{}".format(name, tv)
        df[name2] = l

# Encode text values to indexes(i.e. [1],[2],[3] for red,green,blue).
def encode_text_index(df, name):
    le = preprocessing.LabelEncoder()
    df[name] = le.fit_transform(df[name])
    return le.classes_
```

```

# Encode a numeric column as zscores
def encode_numeric_zscore(df, name, mean=None, sd=None):
    if mean is None:
        mean = df[name].mean()

    if sd is None:
        sd = df[name].std()

    df[name] = (df[name] - mean) / sd

# Convert all missing values in the specified column to the median
def missing_median(df, name):
    med = df[name].median()
    df[name] = df[name].fillna(med)

# Convert all missing values in the specified column to the default
def missing_default(df, name, default_value):
    df[name] = df[name].fillna(default_value)

# Convert a Pandas dataframe to the x,y inputs that TensorFlow needs
def to_xy(df, target):
    result = []
    for x in df.columns:
        if x != target:
            result.append(x)
    # find out the type of the target column. Is it really this hard? :(
    target_type = df[target].dtypes
    target_type = target_type[0] if hasattr(target_type, '__iter__') else
    # Encode to int for classification, float otherwise. TensorFlow likes
    if target_type in (np.int64, np.int32):
        # Classification
        dummies = pd.get_dummies(df[target])
        return df.as_matrix(result).astype(np.float32), dummies.as_matrix()
    else:
        # Regression
        return df.as_matrix(result).astype(np.float32), df.as_matrix([target])

# Nicely formatted time string
def hms_string(sec_elapsed):
    h = int(sec_elapsed / (60 * 60))
    m = int((sec_elapsed % (60 * 60)) / 60)
    s = sec_elapsed % 60
    return "{}:{:>02}:{:>05.2f}".format(h, m, s)

# Regression chart.

```

```

def chart_regression(pred,y,sort=True):
    t = pd.DataFrame({'pred' : pred, 'y' : y.flatten()})
    if sort:
        t.sort_values(by=['y'],inplace=True)
    a = plt.plot(t['y'].tolist(),label='expected')
    b = plt.plot(t['pred'].tolist(),label='prediction')
    plt.ylabel('output')
    plt.legend()
    plt.show()

# Remove all rows where the specified column is +/- sd standard deviations
def remove_outliers(df, name, sd):
    drop_rows = df.index[(np.abs(df[name] - df[name].mean()) >= (sd * df[n
    df.drop(drop_rows, axis=0, inplace=True)

# Encode a column to a range between normalized_low and normalized_high.
def encode_numeric_range(df, name, normalized_low=-1, normalized_high=1,
                          data_low=None, data_high=None):
    if data_low is None:
        data_low = min(df[name])
        data_high = max(df[name])

    df[name] = ((df[name] - data_low) / (data_high - data_low)) \
        * (normalized_high - normalized_low) + normalized_low

```

```

In [4]: # This file is a CSV, just no CSV extension or headers
# Download from: http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html
df = pd.read_csv("/Users/tsaiflow/Downloads/kddcup.data_10_percent_corrected")
# df1 = pd.read_csv("/Users/tsaiflow/Downloads/kddcup.data.corrected", header=0)
print("Read {} rows.".format(len(df)))
#print("Read {} rows.".format(len(df1)))
# df = df.sample(frac=0.1, replace=False) # Uncomment this line to sample
df.dropna(inplace=True,axis=1) # For now, just drop NA's (rows with missing values)
# df1.dropna(inplace=True,axis=1)

# The CSV file has no column heads, so add them
df.columns = [
    'duration',
    'protocol_type',
    'service',#
    'flag',#
    'src_bytes',#
    'dst_bytes',#
    'land',#
    'wrong_fragment',
    'urgent',

    'hot',
    'num_failed_logins',

```

```
'logged_in',
'num_compromised',
'root_shell',
'su_attempted',
'num_root',#
'num_file_creations',
'num_shells',
'num_access_files',
'num_outbound_cmds',
'is_host_login',
'is_guest_login',

'count',
'srv_count',
'serror_rate',
'srv_serror_rate',
'rerror_rate',#
'srv_rerror_rate',
'same_srv_rate',
'diff_srv_rate',#
'srv_diff_host_rate',
'dst_host_count',
'dst_host_srv_count',
'dst_host_same_srv_rate',
'dst_host_diff_srv_rate',#
'dst_host_same_src_port_rate',
'dst_host_srv_diff_host_rate',
'dst_host_serror_rate',
'dst_host_srv_serror_rate',#
'dst_host_rerror_rate',
'dst_host_srv_rerror_rate',

'outcome'
]
```

Read 494021 rows.

```
In [5]: df.head(5)
```

```
Out[5]:
```

	duration	protocol_type	service	flag	src_bytes	dst_bytes	land	wrong_fragment	urgent	ho
0	0	tcp	http	SF	181	5450	0	0	0	(
1	0	tcp	http	SF	239	486	0	0	0	(
2	0	tcp	http	SF	235	1337	0	0	0	(
3	0	tcp	http	SF	219	1337	0	0	0	(
4	0	tcp	http	SF	217	2032	0	0	0	(

5 rows × 42 columns

```
In [6]: df_new = df[['service',
                    'flag',
                    'src_bytes',
                    'dst_bytes',
                    'land',
                    'num_root',
                    'rerror_rate',
                    'diff_srv_rate',
                    'dst_host_diff_srv_rate',
                    'dst_host_srv_serror_rate',
                    'outcome']]
```

```
In [7]: # Now encode the feature vector

#encode_numeric_zscore(df, 'duration')
#encode_text_dummy(df, 'protocol_type')
encode_text_dummy(df_new, 'service')
encode_text_dummy(df_new, 'flag')
encode_numeric_zscore(df_new, 'src_bytes')
encode_numeric_zscore(df_new, 'dst_bytes')
encode_text_dummy(df_new, 'land')
#encode_numeric_zscore(df, 'wrong_fragment')
# encode_numeric_zscore(df, 'urgent')
# encode_numeric_zscore(df, 'hot')
# encode_numeric_zscore(df, 'num_failed_logins')
# encode_text_dummy(df, 'logged_in')
# encode_numeric_zscore(df, 'num_compromised')
# encode_numeric_zscore(df, 'root_shell')
# encode_numeric_zscore(df, 'su_attempted')
encode_numeric_zscore(df_new, 'num_root')
# encode_numeric_zscore(df, 'num_file_creations')
# encode_numeric_zscore(df, 'num_shells')
# encode_numeric_zscore(df, 'num_access_files')
# encode_numeric_zscore(df, 'num_outbound_cmds')
```

```

# encode_text_dummy(df, 'is_host_login')
# encode_text_dummy(df, 'is_guest_login')
# encode_numeric_zscore(df, 'count')
# encode_numeric_zscore(df, 'srv_count')
# encode_numeric_zscore(df, 'serror_rate')
# encode_numeric_zscore(df, 'srv_serror_rate')
encode_numeric_zscore(df_new, 'rerror_rate')
# encode_numeric_zscore(df, 'srv_rerror_rate')
# encode_numeric_zscore(df, 'same_srv_rate')
encode_numeric_zscore(df_new, 'diff_srv_rate')
# encode_numeric_zscore(df, 'srv_diff_host_rate')
# encode_numeric_zscore(df, 'dst_host_count')
# encode_numeric_zscore(df, 'dst_host_srv_count')
# encode_numeric_zscore(df, 'dst_host_same_srv_rate')
encode_numeric_zscore(df_new, 'dst_host_diff_srv_rate')
# encode_numeric_zscore(df, 'dst_host_same_src_port_rate')
# encode_numeric_zscore(df, 'dst_host_srv_diff_host_rate')
# encode_numeric_zscore(df, 'dst_host_serror_rate')
encode_numeric_zscore(df_new, 'dst_host_srv_serror_rate')
# encode_numeric_zscore(df, 'dst_host_rerror_rate')
# encode_numeric_zscore(df, 'dst_host_srv_rerror_rate')
outcomes = encode_text_index(df_new, 'outcome')
num_classes = len(outcomes)

# display 5 rows

df_new.dropna(inplace=True,axis=1)
df_new[0:5]
# This is the numeric feature vector, as it goes to the neural net

```

/anaconda/lib/python3.6/site-packages/ipykernel\_launcher.py:6: Setting  
WithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>  
(<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

Out[7]:

	src_bytes	dst_bytes	num_root	rerror_rate	diff_srv_rate	dst_host_diff_srv_rate	dst_host_srv_s
0	-0.002879	0.138664	-0.00564	-0.24796	-0.255243		-0.282866
1	-0.002820	-0.011578	-0.00564	-0.24796	-0.255243		-0.282866
2	-0.002824	0.014179	-0.00564	-0.24796	-0.255243		-0.282866
3	-0.002840	0.014179	-0.00564	-0.24796	-0.255243		-0.282866

4 -0.002842 0.035214 -0.00564 -0.24796 -0.255243 -0.282866

5 rows × 87 columns

```
In [8]: df1 = pd.read_csv("/Users/tsaiflow/Downloads/kddcup.data.corrected", head
print("Read {} rows.".format(len(df1)))
df1.dropna(inplace=True,axis=1) # For now, just drop NA's (rows with miss
df1.columns = [
    'duration',
    'protocol_type',
    'service',#
    'flag',#
    'src_bytes',#
    'dst_bytes',#
    'land',#
    'wrong_fragment',
    'urgent',

    'hot',
    'num_failed_logins',
    'logged_in',
    'num_compromised',
    'root_shell',
    'su_attempted',
    'num_root',#
    'num_file_creations',
    'num_shells',
    'num_access_files',
    'num_outbound_cmds',
    'is_host_login',
    'is_guest_login',

    'count',
    'srv_count',
    'serror_rate',
    'srv_serror_rate',
    'rerror_rate',#
    'srv_rerror_rate',
    'same_srv_rate',
    'diff_srv_rate',#
    'srv_diff_host_rate',
    'dst_host_count',
    'dst_host_srv_count',
    'dst_host_same_srv_rate',
    'dst_host_diff_srv_rate',#
    'dst_host_same_src_port_rate',
    'dst_host_srv_diff_host_rate',
    'dst_host_serror_rate',
    'dst_host_srv_serror_rate',#
    'dst_host_rerror_rate'
```

```

        'dst_host_srv_error_rate',

        'outcome'
    ]

df1_new = df1[['service',
               'flag',
               'src_bytes',
               'dst_bytes',
               'land',
               'num_root',
               'error_rate',
               'diff_srv_rate',
               'dst_host_diff_srv_rate',
               'dst_host_srv_serror_rate',
               'outcome']]

encode_text_dummy(df1_new, 'service')
encode_text_dummy(df1_new, 'flag')
encode_numeric_zscore(df1_new, 'src_bytes')
encode_numeric_zscore(df1_new, 'dst_bytes')
encode_text_dummy(df1_new, 'land')
encode_numeric_zscore(df1_new, 'num_root')
encode_numeric_zscore(df1_new, 'error_rate')
encode_numeric_zscore(df1_new, 'diff_srv_rate')
encode_numeric_zscore(df1_new, 'dst_host_diff_srv_rate')
encode_numeric_zscore(df1_new, 'dst_host_srv_serror_rate')
outcomes = encode_text_index(df1_new, 'outcome')
# num_classes = len(outcomes)

# display 5 rows

df1_new.dropna(inplace=True,axis=1)

del df1_new['service-aol']
del df1_new['service-harvest']
del df1_new['service-http_2784']
del df1_new['service-http_8001']
#del df1_new['is_host_login-1']
df1_new[0:5]

```

Read 4898431 rows.

/anaconda/lib/python3.6/site-packages/ipykernel\_launcher.py:6: Setting  
WithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas->



docs/stable/indexing.html#indexing-view-versus-copy  
 (http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy)

Out[8]:

	src_bytes	dst_bytes	num_root	error_rate	diff_srv_rate	dst_host_diff_srv_rate	dst_host_srv_s
0	-0.001720	0.068188	-0.003285	-0.248286	-0.256056		-0.282939
1	-0.001777	0.005325	-0.003285	-0.248286	-0.256056		-0.282939
2	-0.001698	0.000208	-0.003285	-0.248286	-0.256056		-0.282939
3	-0.001701	0.001455	-0.003285	-0.248286	-0.256056		-0.282939
4	-0.001695	-0.000942	-0.003285	-0.248286	-0.256056		-0.282939

5 rows × 87 columns

```
In [9]: import numpy
from pandas import read_csv
from sklearn.decomposition import PCA
# load data

array_outcome = df_new[['outcome']].values
array_outcome[array_outcome!=11] = 1
array_outcome[array_outcome==11] = 0

array_outcome_whole = df1_new[['outcome']].values
array_outcome_whole[array_outcome_whole!=11] = 1
array_outcome_whole[array_outcome_whole==11] = 0

del df1_new['outcome']

X_whole = df1_new.values
Y_whole = array_outcome_whole

del df_new['outcome']

X = df_new.values
Y = array_outcome

print(np.shape(X), np.shape(Y))

(494021, 86) (494021, 1)
```

```
In [10]: # #####
# #####PCA#####
# #####

# # feature extraction
# pca = PCA(n_components=20)
# fit = pca.fit(X.T)
# # summarize components
# print(("Explained Variance: %s") % fit.explained_variance_ratio_)
# print(np.shape(fit.components_.T))
# X = fit.components_.T
# # print(fit.components_)
```

```
In [11]: import keras.backend as K
K.clear_session()
```

```
In [12]: # Break into X (predictors) & y (prediction)
# x, y = to_xy(df, 'outcome')
# x_train, y_train = to_xy(df, 'outcome')
# x_test, y_test = to_xy(df1, 'outcome')

# Create a test/train split. 25% test
# Split into train/test
# x_train, x_test, y_train, y_test = train_test_split(
#     X, Y, test_size=0.25, random_state=42)
x_train, x_test, y_train, y_test = train_test_split(
    X, Y, test_size=0.25)
```

```
In [13]: # Create neural net
model = Sequential()

# Used relu for activation function
model.add(Dense(10, input_dim=X.shape[1], kernel_initializer='normal', ac
model.add(Dense(50, input_dim=X.shape[1], kernel_initializer='normal', ac
model.add(Dense(10, input_dim=X.shape[1], kernel_initializer='normal', ac

# output layer
model.add(Dense(1, kernel_initializer='normal'))
model.add(Dense(1, activation = 'sigmoid'))
```

```
In [14]: # TODO: Print the model summary
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 10)	870
dense_2 (Dense)	(None, 50)	550
dense_3 (Dense)	(None, 10)	510
dense_4 (Dense)	(None, 1)	11
dense_5 (Dense)	(None, 1)	2

=====  
 Total params: 1,943  
 Trainable params: 1,943  
 Non-trainable params: 0  
 =====

```
In [15]: class LossHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        # TODO: Create two empty lists, self.loss and self.val_acc
        self.losses = []
        self.accs = []
        self.val_acc = []
        self.val_loss = []
    def on_batch_end(self, batch, logs={}):
        # TODO: This is called at the end of each batch.
        # Add the loss in logs.get('loss') to the loss list
        loss = logs.get('loss')
        acc = logs.get('acc')
        self.losses.append(loss)
        self.accs.append(acc)
    def on_epoch_end(self, epoch, logs):
        # TODO: This is called at the end of each epoch.
        # Add the test accuracy in logs.get('loss') to the val_acc list
        val_acc = logs.get('val_acc')
        val_loss = logs.get('val_loss')
        self.val_acc.append(val_acc)
        self.val_loss.append(val_loss)

    # Create an instance of the history callback
    history_cb = LossHistory()
```

```
In [16]: from keras import optimizers

# optimizer defined
opt = optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-08,
# opt = optimizers.Adam(lr=0.001)

In [17]: ## compile
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['accuracy'])

In [18]: # model.compile(loss='categorical_crossentropy', optimizer='adam', metric
# monitor = EarlyStopping(monitor='val_loss', min_delta=1e-3, patience=5,

In [19]: # model.fit(x,y,validation_data=(x_test,y_test),callbacks=[monitor],verbo
batch_size = 100
epochs = 5
model.fit(x_train,y_train,callbacks=[history_cb],verbose=1,epochs=epochs,
          validation_data=(x_test,y_test))

Train on 370515 samples, validate on 123506 samples
Epoch 1/5
370515/370515 [=====] - 15s - loss: 0.0580 -
acc: 0.9841 - val_loss: 0.0239 - val_acc: 0.9913
Epoch 2/5
370515/370515 [=====] - 15s - loss: 0.0219 -
acc: 0.9946 - val_loss: 0.0195 - val_acc: 0.9949
Epoch 3/5
370515/370515 [=====] - 15s - loss: 0.0198 -
acc: 0.9950 - val_loss: 0.0189 - val_acc: 0.9949
Epoch 4/5
370515/370515 [=====] - 15s - loss: 0.0188 -
acc: 0.9950 - val_loss: 0.0174 - val_acc: 0.9958
Epoch 5/5
370515/370515 [=====] - 14s - loss: 0.0179 -
acc: 0.9954 - val_loss: 0.0180 - val_acc: 0.9947


Out[19]: <keras.callbacks.History at 0x19735cdd8>
```

```
In [20]: yhat = model.predict(x_test)
yhat[yhat>0.5] = 1
yhat[yhat<=0.5] = 0
count = 0
yhat_size = np.shape(yhat)[0]

for i in range (yhat_size):
    if yhat[i] != y_test[i]:
        count = count + 1
print("Accuracy is: %s" %(1-count/yhat_size))
print("# of wrongly classified: %s" %(count))
```


Accuracy is: 0.9946642268391819  
# of wrongly classified: 659

```
In [21]: valAcc = history_cb.val_acc
valLoss = history_cb.val_loss
epoch_it = np.arange(1,6)
plt.subplot(121)
plt.plot(epoch_it, valAcc)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Validation accuracy')
plt.tight_layout()
plt.subplot(122)
plt.plot(epoch_it, valLoss)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Validation loss')
plt.tight_layout()
plt.show()
```




```
In [22]: Ls = history_cb.losses
Acc = history_cb.accs
ntr = np.shape(x_train)[0]
epochNum = []
for i in range(18530):
    epochNum.append(i*100/ntr)

matplotlib.pyplot.semilogy(epochNum, Ls)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.tight_layout()
plt.show()
```



```
In [23]: matplotlib.pyplot.semilogy(epochNum, Acc)
# plt.plot(epochNum, Acc)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.tight_layout()
# plt.axis([0.3,3,0.96,1])
plt.show()
```



```
In [24]: #####
#####details#####
#####
##The following is the trained model validated on the whole dataset
```

```
In [25]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [26]: y_pred_whole = model.predict_classes(X_whole, verbose=1)
# print(y_pred[:5])
```

```
4896480/4898431 [=====>.] - ETA: 0s
```

```
In [27]: count = 0
yhat_size = np.shape(y_pred_whole)[0]

for i in range (yhat_size):
    if y_pred_whole[i] != Y_whole[i]:
        count = count + 1
print("Accuracy is: %s" %(1-count/yhat_size))
print("# of wrongly classified: %s" %(count))
```

```
Accuracy is: 0.9977035503817446
# of wrongly classified: 11249
```

```
In [28]: print(classification_report(Y_whole, y_pred_whole))
```

	precision	recall	f1-score	support
0	1.00	0.99	0.99	972781
1	1.00	1.00	1.00	3925650
avg / total	1.00	1.00	1.00	4898431

```
In [29]: Cm = confusion_matrix(Y_whole,y_pred_whole)
C = np.sum(Cm)
Cm = Cm/C
print('Confusion Matrix:')
print(np.array_str(Cm, precision=4, suppress_small=True))
```

```
Confusion Matrix:
[[ 0.1969  0.0017]
 [ 0.0006  0.8008]]
```