

```
In [1]: # Imports for this Notebook
import pandas as pd
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn import metrics
from keras.models import Sequential
from keras.layers.core import Dense, Activation
from keras.callbacks import EarlyStopping
```

Using TensorFlow backend.

```
In [2]: from sklearn import preprocessing
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import shutil
import os
import keras
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: # Encode text values to dummy variables(i.e. [1,0,0],[0,1,0],[0,0,1] for
def encode_text_dummy(df, name):
    dummies = pd.get_dummies(df[name])
    for x in dummies.columns:
        dummy_name = "{}-{}".format(name, x)
        df[dummy_name] = dummies[x]
    df.drop(name, axis=1, inplace=True)

# Encode text values to a single dummy variable. The new columns (which
# at every location where the original column (name) matches each of the
# each target value.
def encode_text_single_dummy(df, name, target_values):
    for tv in target_values:
        l = list(df[name].astype(str))
        l = [1 if str(x) == str(tv) else 0 for x in l]
        name2 = "{}-{}".format(name, tv)
        df[name2] = l

# Encode text values to indexes(i.e. [1],[2],[3] for red,green,blue).
def encode_text_index(df, name):
    le = preprocessing.LabelEncoder()
    df[name] = le.fit_transform(df[name])
    return le.classes_
```

```

# Encode a numeric column as zscores
def encode_numeric_zscore(df, name, mean=None, sd=None):
    if mean is None:
        mean = df[name].mean()

    if sd is None:
        sd = df[name].std()

    df[name] = (df[name] - mean) / sd

# Convert all missing values in the specified column to the median
def missing_median(df, name):
    med = df[name].median()
    df[name] = df[name].fillna(med)

# Convert all missing values in the specified column to the default
def missing_default(df, name, default_value):
    df[name] = df[name].fillna(default_value)

# Convert a Pandas dataframe to the x,y inputs that TensorFlow needs
def to_xy(df, target):
    result = []
    for x in df.columns:
        if x != target:
            result.append(x)
    # find out the type of the target column. Is it really this hard? :(
    target_type = df[target].dtypes
    target_type = target_type[0] if hasattr(target_type, '__iter__') else
    # Encode to int for classification, float otherwise. TensorFlow likes
    if target_type in (np.int64, np.int32):
        # Classification
        dummies = pd.get_dummies(df[target])
        return df.as_matrix(result).astype(np.float32), dummies.as_matrix
    else:
        # Regression
        return df.as_matrix(result).astype(np.float32), df.as_matrix([tar

# Nicely formatted time string
def hms_string(sec_elapsed):
    h = int(sec_elapsed / (60 * 60))
    m = int((sec_elapsed % (60 * 60)) / 60)
    s = sec_elapsed % 60
    return "{}:{:>02}:{:>05.2f}".format(h, m, s)

# Regression chart.

```

```

def chart_regression(pred,y,sort=True):
    t = pd.DataFrame({'pred' : pred, 'y' : y.flatten()})
    if sort:
        t.sort_values(by=['y'],inplace=True)
    a = plt.plot(t['y'].tolist(),label='expected')
    b = plt.plot(t['pred'].tolist(),label='prediction')
    plt.ylabel('output')
    plt.legend()
    plt.show()

# Remove all rows where the specified column is +/- sd standard deviation
def remove_outliers(df, name, sd):
    drop_rows = df.index[(np.abs(df[name] - df[name].mean()) >= (sd * df[
df.drop(drop_rows, axis=0, inplace=True)

# Encode a column to a range between normalized_low and normalized_high.
def encode_numeric_range(df, name, normalized_low=-1, normalized_high=1,
                        data_low=None, data_high=None):
    if data_low is None:
        data_low = min(df[name])
        data_high = max(df[name])

    df[name] = ((df[name] - data_low) / (data_high - data_low)) \
        * (normalized_high - normalized_low) + normalized_low

```

In [4]:

```

# This file is a CSV, just no CSV extension or headers
# Download from: http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html
df = pd.read_csv("/Users/tsaiflow/Downloads/kddcup.data_10_percent_correc

print("Read {} rows.".format(len(df)))
#print("Read {} rows.".format(len(df1)))
# df = df.sample(frac=0.1, replace=False) # Uncomment this line to sample
df.dropna(inplace=True,axis=1) # For now, just drop NA's (rows with missi
# df1.dropna(inplace=True,axis=1)

# The CSV file has no column heads, so add them
df.columns = [
    'duration',
    'protocol_type',
    'service',
    'flag',
    'src_bytes',
    'dst_bytes',
    'land',
    'wrong_fragment',
    'urgent',

    'hot',

```

```
'num_failed_logins',  
'logged_in',  
'num_compromised',  
'root_shell',  
'su_attempted',  
'num_root',  
'num_file_creations',  
'num_shells',  
'num_access_files',  
'num_outbound_cmds',  
'is_host_login',  
'is_guest_login',  
  
'count',  
'srv_count',  
'serror_rate',  
'srv_serror_rate',  
'rerror_rate',  
'srv_rerror_rate',  
'same_srv_rate',  
'diff_srv_rate',  
'srv_diff_host_rate',  
'dst_host_count',  
'dst_host_srv_count',  
'dst_host_same_srv_rate',  
'dst_host_diff_srv_rate',  
'dst_host_same_src_port_rate',  
'dst_host_srv_diff_host_rate',  
'dst_host_serror_rate',  
'dst_host_srv_serror_rate',  
'dst_host_rerror_rate',  
'dst_host_srv_rerror_rate',  
  
'outcome'  
]
```

Read 494021 rows.

In [5]: `df.head(5)`

Out[5]:

|   | duration | protocol_type | service | flag | src_bytes | dst_bytes | land | wrong_fragment | urgent | ho |
|---|----------|---------------|---------|------|-----------|-----------|------|----------------|--------|----|
| 0 | 0        | tcp           | http    | SF   | 181       | 5450      | 0    | 0              | 0      | (  |
| 1 | 0        | tcp           | http    | SF   | 239       | 486       | 0    | 0              | 0      | (  |
| 2 | 0        | tcp           | http    | SF   | 235       | 1337      | 0    | 0              | 0      | (  |
| 3 | 0        | tcp           | http    | SF   | 219       | 1337      | 0    | 0              | 0      | (  |
| 4 | 0        | tcp           | http    | SF   | 217       | 2032      | 0    | 0              | 0      | (  |

5 rows × 42 columns

```
In [6]: # Now encode the feature vector

encode_numeric_zscore(df, 'duration')
encode_text_dummy(df, 'protocol_type')
encode_text_dummy(df, 'service')
encode_text_dummy(df, 'flag')
encode_numeric_zscore(df, 'src_bytes')
encode_numeric_zscore(df, 'dst_bytes')
encode_text_dummy(df, 'land')
encode_numeric_zscore(df, 'wrong_fragment')
encode_numeric_zscore(df, 'urgent')
encode_numeric_zscore(df, 'hot')
encode_numeric_zscore(df, 'num_failed_logins')
encode_text_dummy(df, 'logged_in')
encode_numeric_zscore(df, 'num_compromised')
encode_numeric_zscore(df, 'root_shell')
encode_numeric_zscore(df, 'su_attempted')
encode_numeric_zscore(df, 'num_root')
encode_numeric_zscore(df, 'num_file_creations')
encode_numeric_zscore(df, 'num_shells')
encode_numeric_zscore(df, 'num_access_files')
encode_numeric_zscore(df, 'num_outbound_cmds')
encode_text_dummy(df, 'is_host_login')
encode_text_dummy(df, 'is_guest_login')
encode_numeric_zscore(df, 'count')
encode_numeric_zscore(df, 'srv_count')
encode_numeric_zscore(df, 'serror_rate')
encode_numeric_zscore(df, 'srv_serror_rate')
encode_numeric_zscore(df, 'rerror_rate')
encode_numeric_zscore(df, 'srv_rerror_rate')
encode_numeric_zscore(df, 'same_srv_rate')
encode_numeric_zscore(df, 'diff_srv_rate')
encode_numeric_zscore(df, 'srv_diff_host_rate')
encode_numeric_zscore(df, 'dst_host_count')
```

```

encode_numeric_zscore(df, 'dst_host_srv_count')
encode_numeric_zscore(df, 'dst_host_same_srv_rate')
encode_numeric_zscore(df, 'dst_host_diff_srv_rate')
encode_numeric_zscore(df, 'dst_host_same_src_port_rate')
encode_numeric_zscore(df, 'dst_host_srv_diff_host_rate')
encode_numeric_zscore(df, 'dst_host_serror_rate')
encode_numeric_zscore(df, 'dst_host_srv_serror_rate')
encode_numeric_zscore(df, 'dst_host_rerror_rate')
encode_numeric_zscore(df, 'dst_host_srv_rerror_rate')
outcomes = encode_text_index(df, 'outcome')
num_classes = len(outcomes)

# display 5 rows

df.dropna(inplace=True,axis=1)
df[0:5]
# This is the numeric feature vector, as it goes to the neural net

```

Out[6]:

|   | duration  | src_bytes | dst_bytes | wrong_fragment | urgent    | hot       | num_failed_logins | num_ |
|---|-----------|-----------|-----------|----------------|-----------|-----------|-------------------|------|
| 0 | -0.067792 | -0.002879 | 0.138664  | -0.04772       | -0.002571 | -0.044136 | -0.009782         |      |
| 1 | -0.067792 | -0.002820 | -0.011578 | -0.04772       | -0.002571 | -0.044136 | -0.009782         |      |
| 2 | -0.067792 | -0.002824 | 0.014179  | -0.04772       | -0.002571 | -0.044136 | -0.009782         |      |
| 3 | -0.067792 | -0.002840 | 0.014179  | -0.04772       | -0.002571 | -0.044136 | -0.009782         |      |
| 4 | -0.067792 | -0.002842 | 0.035214  | -0.04772       | -0.002571 | -0.044136 | -0.009782         |      |

5 rows × 121 columns

```

In [7]: df1 = pd.read_csv("/Users/tsaiflow/Downloads/kddcup.data.corrected", header=0)
print("Read {} rows.".format(len(df1)))
# print("Read {} rows.".format(len(df1)))
# df = df.sample(frac=0.1, replace=False) # Uncomment this line to sample
df1.dropna(inplace=True,axis=1) # For now, just drop NA's (rows with missing values)
# df1.dropna(inplace=True,axis=1)

# The CSV file has no column heads, so add them
df1.columns = [
    'duration',
    'protocol_type',
    'service',
    'flag',
    'src_bytes',
    'dst_bytes',
    'land',
    'wrong_fragment',
    'urgent',

```

```

    'hot',
    'num_failed_logins',
    'logged_in',
    'num_compromised',
    'root_shell',
    'su_attempted',
    'num_root',
    'num_file_creations',
    'num_shells',
    'num_access_files',
    'num_outbound_cmds',
    'is_host_login',
    'is_guest_login',

    'count',
    'srv_count',
    'serror_rate',
    'srv_serror_rate',
    'rerror_rate',
    'srv_rerror_rate',
    'same_srv_rate',
    'diff_srv_rate',
    'srv_diff_host_rate',
    'dst_host_count',
    'dst_host_srv_count',
    'dst_host_same_srv_rate',
    'dst_host_diff_srv_rate',
    'dst_host_same_src_port_rate',
    'dst_host_srv_diff_host_rate',
    'dst_host_serror_rate',
    'dst_host_srv_serror_rate',
    'dst_host_rerror_rate',
    'dst_host_srv_rerror_rate',

    'outcome'
]

# Now encode the feature vector

encode_numeric_zscore(df1, 'duration')
encode_text_dummy(df1, 'protocol_type')
encode_text_dummy(df1, 'service')
encode_text_dummy(df1, 'flag')
encode_numeric_zscore(df1, 'src_bytes')
encode_numeric_zscore(df1, 'dst_bytes')
encode_text_dummy(df1, 'land')
encode_numeric_zscore(df1, 'wrong_fragment')
encode_numeric_zscore(df1, 'urgent')
encode_numeric_zscore(df1, 'hot')
encode_numeric_zscore(df1, 'num_failed_logins')

```

```

-----\-----, -----\-----,
encode_text_dummy(df1, 'logged_in')
encode_numeric_zscore(df1, 'num_compromised')
encode_numeric_zscore(df1, 'root_shell')
encode_numeric_zscore(df1, 'su_attempted')
encode_numeric_zscore(df1, 'num_root')
encode_numeric_zscore(df1, 'num_file_creations')
encode_numeric_zscore(df1, 'num_shells')
encode_numeric_zscore(df1, 'num_access_files')
encode_numeric_zscore(df1, 'num_outbound_cmds')
encode_text_dummy(df1, 'is_host_login')
encode_text_dummy(df1, 'is_guest_login')
encode_numeric_zscore(df1, 'count')
encode_numeric_zscore(df1, 'srv_count')
encode_numeric_zscore(df1, 'serror_rate')
encode_numeric_zscore(df1, 'srv_serror_rate')
encode_numeric_zscore(df1, 'rerror_rate')
encode_numeric_zscore(df1, 'srv_rerror_rate')
encode_numeric_zscore(df1, 'same_srv_rate')
encode_numeric_zscore(df1, 'diff_srv_rate')
encode_numeric_zscore(df1, 'srv_diff_host_rate')
encode_numeric_zscore(df1, 'dst_host_count')
encode_numeric_zscore(df1, 'dst_host_srv_count')
encode_numeric_zscore(df1, 'dst_host_same_srv_rate')
encode_numeric_zscore(df1, 'dst_host_diff_srv_rate')
encode_numeric_zscore(df1, 'dst_host_same_src_port_rate')
encode_numeric_zscore(df1, 'dst_host_srv_diff_host_rate')
encode_numeric_zscore(df1, 'dst_host_serror_rate')
encode_numeric_zscore(df1, 'dst_host_srv_serror_rate')
encode_numeric_zscore(df1, 'dst_host_rerror_rate')
encode_numeric_zscore(df1, 'dst_host_srv_rerror_rate')
outcomes = encode_text_index(df1, 'outcome')
#num_classes = len(outcomes)

# display 5 rows

df1.dropna(inplace=True,axis=1)
del df1['service-aol']
del df1['service-harvest']
del df1['service-http_2784']
del df1['service-http_8001']
del df1['is_host_login-1']
df1[0:5]
# This is the numeric feature vector, as it goes to the neural net

Read 4898431 rows.

```

Out[7]:

|   | duration  | src_bytes | dst_bytes | wrong_fragment | urgent    | hot       | num_failed_logins | num_ |
|---|-----------|-----------|-----------|----------------|-----------|-----------|-------------------|------|
| 0 | -0.066833 | -0.001720 | 0.068188  | -0.015139      | -0.001103 | -0.026521 | -0.004391         |      |



|   |           |           |           |           |           |           |           |
|---|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 1 | -0.066833 | -0.001777 | 0.005325  | -0.015139 | -0.001103 | -0.026521 | -0.004391 |
| 2 | -0.066833 | -0.001698 | 0.000208  | -0.015139 | -0.001103 | -0.026521 | -0.004391 |
| 3 | -0.066833 | -0.001701 | 0.001455  | -0.015139 | -0.001103 | -0.026521 | -0.004391 |
| 4 | -0.066833 | -0.001695 | -0.000942 | -0.015139 | -0.001103 | -0.026521 | -0.004391 |

5 rows × 121 columns

```
In [8]: import numpy
from pandas import read_csv
from sklearn.decomposition import PCA
# load data

array_outcome_whole = df1[['outcome']].values
array_outcome_whole[array_outcome_whole!=11] = 1
array_outcome_whole[array_outcome_whole==11] = 0

del df1['outcome']

X_whole = df1.values
Y_whole = array_outcome_whole
```

```
In [9]: array_outcome = df[['outcome']].values
array_outcome[array_outcome!=11] = 1
array_outcome[array_outcome==11] = 0

del df['outcome']

X = df.values
Y = array_outcome
print(np.shape(X), np.shape(Y))

(494021, 120) (494021, 1)
```

```

In [10]: # #####
# #####PCA#####
# #####

# # feature extraction
pca = PCA(n_components=20)
fit = pca.fit(X.T)

# summarize components
print(("Explained Variance: %s") % fit.explained_variance_ratio_)
print(np.shape(fit.components_.T))
X = fit.components_.T
# print(fit.components_)

Explained Variance: [ 0.22845019  0.12651505  0.11281686  0.07454234
 0.06876901  0.03971399
 0.02952583  0.02902571  0.02674098  0.02584283  0.02534576  0.025217
 73
 0.02452765  0.02296654  0.02176682  0.02160914  0.01961439  0.018340
 38
 0.01812486  0.00973112]
(494021, 20)

In [11]: fit_whole = pca.fit(X_whole.T)
print(("Explained Variance: %s") % fit_whole.explained_variance_ratio_)
print(np.shape(fit_whole.components_.T))
X_whole = fit_whole.components_.T

Explained Variance: [ 0.22962152  0.1264887   0.11289027  0.07114857
 0.05875115  0.03964593
 0.03134808  0.02642499  0.02582551  0.02562436  0.0254813   0.025451
 3
 0.02535263  0.0249264   0.02399793  0.02320513  0.02133873  0.020965
 31
 0.018345    0.012523   ]
(4898431, 20)

In [12]: import keras.backend as K
K.clear_session()

In [13]: # Break into X (predictors) & y (prediction)
# Create a test/train split. 25% test
# Split into train/test

x_train, x_test, y_train, y_test = train_test_split(
    X, Y, test_size=0.25)

```

```
In [14]: # Create neural net
model = Sequential()

# Used relu for activation function
model.add(Dense(10, input_dim=X.shape[1], kernel_initializer='normal', act
model.add(Dense(50, input_dim=X.shape[1], kernel_initializer='normal', act
model.add(Dense(10, input_dim=X.shape[1], kernel_initializer='normal', act

# output layer
model.add(Dense(1, kernel_initializer='normal'))
# model.add(Dense(y.shape[1],activation='softmax'))
model.add(Dense(1,activation = 'sigmoid'))
```

```
In [15]: # TODO: Print the model summary
model.summary()
```

| Layer (type)            | Output Shape | Param # |
|-------------------------|--------------|---------|
| =====                   |              |         |
| dense_1 (Dense)         | (None, 10)   | 210     |
| dense_2 (Dense)         | (None, 50)   | 550     |
| dense_3 (Dense)         | (None, 10)   | 510     |
| dense_4 (Dense)         | (None, 1)    | 11      |
| dense_5 (Dense)         | (None, 1)    | 2       |
| =====                   |              |         |
| Total params: 1,283     |              |         |
| Trainable params: 1,283 |              |         |
| Non-trainable params: 0 |              |         |
| =====                   |              |         |

```
In [16]: class LossHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        # TODO: Create two empty lists, self.loss and self.val_acc
        self.losses = []
        self.accs = []
        self.val_acc = []
        self.val_loss = []
    def on_batch_end(self, batch, logs={}):
        # TODO: This is called at the end of each batch.
        # Add the loss in logs.get('loss') to the loss list
        loss = logs.get('loss')
        acc = logs.get('acc')
        self.losses.append(loss)
        self.accs.append(acc)
    def on_epoch_end(self, epoch, logs):
        # TODO: This is called at the end of each epoch.
        # Add the test accuracy in logs.get('loss') to the val_acc list
        val_acc = logs.get('val_acc')
        val_loss = logs.get('val_loss')
        self.val_acc.append(val_acc)
        self.val_loss.append(val_loss)

# Create an instance of the history callback
history_cb = LossHistory()
```

```
In [17]: from keras import optimizers

# optimizer defined
opt = optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-08,
```

```
In [18]: ## compile model
model.compile(optimizer=opt,
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
In [19]: batch_size = 100
epochs = 5
model.fit(x_train,y_train,callbacks=[history_cb],validation_data=(x_test,
```

Train on 370515 samples, validate on 123506 samples

Epoch 1/5

370515/370515 [=====] - 12s - loss: 0.2653 -

acc: 0.8971 - val\_loss: 0.0405 - val\_acc: 0.9891

Epoch 2/5

370515/370515 [=====] - 12s - loss: 0.0395 -

acc: 0.9903 - val\_loss: 0.0369 - val\_acc: 0.9915

Epoch 3/5

370515/370515 [=====] - 12s - loss: 0.0376 -

acc: 0.9910 - val\_loss: 0.0362 - val\_acc: 0.9915

Epoch 4/5

370515/370515 [=====] - 12s - loss: 0.0369 -

acc: 0.9912 - val\_loss: 0.0356 - val\_acc: 0.9916

Epoch 5/5

370515/370515 [=====] - 12s - loss: 0.0364 -

acc: 0.9913 - val\_loss: 0.0351 - val\_acc: 0.9917

Out[19]: <keras.callbacks.History at 0x121d8de48>


```
In [20]: yhat = model.predict(x_test)
yhat[yhat>0.5] = 1
yhat[yhat<=0.5] = 0
count = 0
yhat_size = np.shape(yhat)[0]

for i in range (yhat_size):
    if yhat[i] != y_test[i]:
        count = count + 1
print("Accuracy is: %s" %(1-count/yhat_size))
print("# of wrongly classified: %s" %(count))
```

Accuracy is: 0.9917089048305345


# of wrongly classified: 1024

```
In [21]: valAcc = history_cb.val_acc
valLoss = history_cb.val_loss
epoch_it = np.arange(1,6)
plt.subplot(121)
plt.plot(epoch_it, valAcc)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Validation accuracy')
plt.tight_layout()
plt.subplot(122)
plt.plot(epoch_it, valLoss)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Validation loss')
plt.tight_layout()
plt.show()
```




```
In [22]: Ls = history_cb.losses
Acc = history_cb.accs
ntr = np.shape(x_train)[0]
epochNum = []
for i in range(18530):
    epochNum.append(i*100/ntr)

matplotlib.pyplot.semilogy(epochNum, Ls)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.tight_layout()
plt.show()
```



```
In [23]: matplotlib.pyplot.semilogy(epochNum, Acc)
# plt.plot(epochNum, Acc)
plt.grid()
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.tight_layout()
# plt.axis([0.3,3,0.96,1])
plt.show()
```



```
In [24]: #####
#####details#####
#####
```

```
In [25]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [26]: y_pred_whole = model.predict_classes(X_whole, verbose=1)
#print(y_pred_whole[:5])
```

```
4894816/4898431 [=====>.] - ETA: 0s
```



```
In [27]: # yhat[yhat>0.5] = 1
# yhat[yhat<=0.5] = 0
count = 0
yhat_size = np.shape(y_pred_whole)[0]

for i in range (yhat_size):
    if y_pred_whole[i] != Y_whole[i]:
        count = count + 1
print("Accuracy is: %s" %(1-count/yhat_size))
print("# of wrongly classified: %s" %(count))
```

```
Accuracy is: 0.8005046922167527
# of wrongly classified: 977214
```

```
In [28]: print(classification_report(Y_whole, y_pred_whole))
```

|             | precision | recall | f1-score | support |
|-------------|-----------|--------|----------|---------|
| 0           | 0.45      | 0.02   | 0.04     | 972781  |
| 1           | 0.80      | 0.99   | 0.89     | 3925650 |
| avg / total | 0.73      | 0.80   | 0.72     | 4898431 |

```
In [29]: Cm = confusion_matrix(Y_whole,y_pred_whole)
C = np.sum(Cm)
Cm = Cm/C
print('Confusion Matrix:')
print(np.array_str(Cm, precision=4, suppress_small=True))
```

```
Confusion Matrix:
[[ 0.0042  0.1944]
 [ 0.0051  0.7963]]
```

```
In [ ]:
```